

Verification Goes Where the Agent Is Already Looking: Intent-Aligned Triage of Inherited Memory Under Budget

Kazuki Nakayashiki
Glasp

Abstract

A long-horizon agent inherits organizational memory it did not write. Giving every record a provenance link does not make that memory safe: an agent holding hundreds of inherited beliefs can follow only a few links before it acts, so the reliability question is not whether provenance exists but which links get followed. We put six production models in a controlled scenario where one of six inherited memories has lost a true negative caveat and at most k source records can be pulled before committing, and report five pre-registered experiments.

Where scarce verification goes. Allocation tracks the agent’s current plan. Across 1020 episodes the corrupted memory is verified in 236/236 episodes whose first-pass intent it backs and 464/784 otherwise, with no counterexample; moving the same corruption onto the intended path makes it caught 75/75. We report this as an allocation *signature* rather than a causal claim: intent and lookups are named in the same response.

Whether it matters later. It does, causally. Randomizing which source records are carried into a later decision, after the budget is spent and the situation has shifted, moves the corrupted-direction rate from 139/150 to 3/150 and takes unguarded commitment from 39/150 to 0/150. A replay stripping the steering instruction and the conversation history agrees within a point.

What the threat model actually is. Two boundary results narrow it. With body length and surface hedging matched, a stated caveat *suppresses* verification and a memory hedging about something irrelevant is checked more often than one whose material caveat was silently deleted — silence is stealthier than qualification, and our earlier reading of that contrast as omission detection was wrong. And the benign consolidation chain we tested does not produce the corruption our own experiments install: over 6 generations quantified negatives mostly survive (83%) while scope (61%) and prohibition erode, and a body with no negative content left appears in 0/360 chain-generations. In the consolidation setup we test, the failure that appears instead is a memory that stays factually accurate while losing the limits that made it safe to act on.

All hypotheses and scoring rules were committed before the corresponding model calls; every number reproduces from released episode files with no API access.

1 Introduction

An agent that runs for weeks does not act from raw evidence. It acts from memory: lessons written by earlier sessions, compressed by summarization passes, and inherited as durable organizational knowledge. Each generation of compression is lossy, and the usual worry is that what it drops is exactly what constrains action: “Discounting lifted signups 31% and revenue 18% *but destroyed retention*” compressing to “discounting lifted signups 31% and revenue 18%.” Nothing false has been written; a constraint has silently stopped being carried forward.

We tested that premise rather than assuming it, and it does not survive (§8). Running a real consolidation chain for 6 generations, the quantified negative result is close to the last thing

compression discards — still stated in 83% of chains when lessons are cut to seventy characters. What erodes is the *scope* of a finding (61% survive) and its normative force: the sentence saying not to do it anyway, and a body carrying no negative content at all appears in 0/360 chain-generations.

This matters for how the rest of the paper should be read. The corruption we install by hand — a lesson stripped of its negative content — is a real failure mode and the one that isolates the allocation question cleanly, but it is not what the benign consolidation process we tested produces. Experiments 1–4 establish what an agent does when such a belief is present; they are not evidence that summarisation routinely creates one.

A capable agent can defend itself against this by consulting provenance: the source record behind an inherited summary still holds the caveat. Prior work has shown that models will follow provenance chains when a belief looks suspicious, and a growing literature studies how to spend a limited verification budget on candidate *outputs*. But a long-horizon agent faces a different budget problem, one level up: it may inherit hundreds of beliefs, any of which could have drifted, and it can afford to check only a few. The reliability question is not whether the agent *can* verify — it is *what it chooses to verify*.

Stated in retrieval terms, this is a selective-retrieval problem over an agent’s own memory store: given a set of candidate records, a budget of k provenance lookups, and no labels indicating which record has drifted, which records does the agent choose to retrieve? What is being ranked is not relevance to a user query but *worth checking* for the decision at hand, and the ground truth — which record lost its caveat — is known to us because we removed it.

This paper measures that choice directly. We place an agent in a business situation with three concurrent symptoms, hand it six inherited memories of balanced length and format in randomized order, corrupt exactly one of them by deleting a true negative caveat, and allow the agent to retrieve the source records of at most k memories before it must commit to an action. Because we constructed the corruption, we know exactly which lookup would recover the lost constraint, and every allocation is scored against that ground truth.

The design has a property we consider essential: the corrupted belief is *not* the one the situation makes salient. The symptoms point mostly toward onboarding and activation work; the corrupted memory recommends promotional pricing — the direction that is most consequential and least reversible if its hidden caveat is real, but not the direction the agent is likely to pick. A verification policy that follows risk should check it. A verification policy that follows intent will check something else.

Our findings, from seventeen hypotheses pre-registered across five experiments, each committed before that experiment’s first model call:

- **Allocation follows intent.** Across all drifted episodes at budget two, 65% of verification credits go to memories backing the model’s own first-pass intended action (uniform-allocation expectation 30%). The corrupted belief is verified in 48/48 of episodes where the model already intends pricing, against 136/252 where it does not (Cochran–Mantel–Haenszel, stratified by model, $p = 2.4 \times 10^{-6}$).
- **Verification shifts with the caveat, unevenly.** The same memory attracts verification 184/300 of the time when its caveat is missing versus 38/150 when present ($p = 2.6 \times 10^{-12}$) — but per-model rates at the same task range from near-always to near-never.
- **Symmetry confirms the mechanism.** When we instead corrupt the intent-aligned memory (onboarding), it is verified 75/75 of the time, against 92/150 for the intent-misaligned target in the matched arm ($p = 1.4 \times 10^{-12}$). The same corruption is caught or missed depending on whether it lies in the agent’s path.

- **Instructions help but do not equalize.** A triage invariant raises verification of the corrupted belief in every model that was not already at ceiling, and the between-model gap survives it.
- **The cost arrives when the context shifts.** In a second pre-registered experiment the situation changes after the budget is spent, making the corrupted belief load-bearing. The pre-registered strict failure metrics return zero — self-reported hedging saturates — but exploratorily, phase-two behavior splits on what the earlier budget recovered: 6/94 corrupted-direction choices with the caveat in context versus 44/56 without, the same direction in every model, and the triage instruction (randomized) halves the rate (50/150 $\rightarrow$ 21/150). The study’s only non-guarded moves into the corrupted direction occur under drift (9 vs. 0 clean).
- **Stated judgment and allocation dissociate.** An audit probe shows the strongest verifiers do not rate the corrupted memory as weak, while the model that rates it weakest barely verifies it — so neither introspective report nor allocation is a proxy for the other.
- **What the allocation recovers causally changes later behaviour.** Randomizing which source records are carried into the later decision — so nothing is self-selected — moves the corrupted-direction rate from 139/150 to 3/150, and removes unguarded commitment to it entirely (39/150 against 0/150). This identifies the downstream half of the chain; it does not identify the upstream half, since intent and lookups are named in the same response. We report an allocation *signature*, not a demonstration that intent causes it (§6).
- **The mechanism is not omission detection.** With body length and surface hedging matched, the true caveat *suppresses* verification (25%), and a memory hedging about something irrelevant (64%) is checked more often than one whose material caveat was deleted (49%). Silence is stealthier than qualification (§7).
- **The consolidation chain we tested does not produce this corruption.** Over 6 generations, quantified negatives mostly survive (83%) while scope restrictions (61%) and prohibitions erode, and a body with no negative content left appears in 0/360. Within this setup the drift keeps the fact and loses the limit: a memory can stay accurate while becoming decision-dangerous (§8).
- **No strict behavioral failures.** Under the pre-registered harmful-decision rule, 0 episodes of 1020 in Experiment 1 and 0 of 450 in Experiment 2 end badly. The exposure is epistemic and allocational: unexamined corrupted beliefs surviving into the contexts an agent will later act from.

2 Related Work

Budgeted verification of model outputs. A recent line of work allocates limited verification compute across candidate outputs or reasoning steps: belief-guided escalation of high-risk outputs [Yuan et al., 2026a], budget control for search agents [Fang et al., 2026], and analyses of when allocation policies break down under heterogeneous signal quality [Yang, 2026]. Benchmarks probe whether agents manage tool budgets at all [Wang and Xu, 2026, Lin et al., 2026]. This literature allocates verification over *candidate answers or actions*; we allocate over *inherited beliefs*, and measure the direction of the allocation rather than its total.

Memory integrity and poisoning. Memory poisoning studies adversarial insertion of content into agent memory [Dash et al., 2026, Xie et al., 2026], and defenses increasingly use provenance — trusting sources rather than content [Singh, 2026]. Our corruption is not adversarial: nothing false is inserted and the remaining sentence is true. It is not, however, what the benign consolidation chain we test in §8 produces: that chain erodes scope and prohibition while keeping quantified negatives, and never strips a lesson of negative content entirely. Three sources are worth keeping apart — the constraint loss the chain we tested does produce; stronger operational transformations we did not test, such as aggressive rewriting, lossy migration or schema conversion; and adversarial poisoning, a different threat model. Our operator isolates the allocation question cleanly and belongs to the second group, not the first. The threat model is operational, not attacker-driven, extending the fault-injection tradition [Tan et al., 2026] from infrastructure and transport faults to the semantic layer.

Provenance and self-auditing. Execution-provenance surveys catalogue how evidence, memory, and actions connect in agent traces [Wang et al., 2026]; self-auditing repairs unsupported intermediate beliefs before they commit [Yuan et al., 2026b]; and provenance-sensitivity audits ask whether action selection responds to source information at all [Liao, 2026]. These works assume the audit target is given or audit everything. We study the upstream selection problem: with hundreds of potential targets and a budget of two, the choice of *what* to audit is itself the reliability variable.

Consolidation and source authority. Concurrent work reports the same asymmetry we find in §8 from a different direction. Zhan et al. [2026] report a closely related authority-preservation failure: consolidation can preserve a claim while losing the constraints on how that claim may be used, observed in 48 of 49 configurations across seven agent-memory consolidators and seven backbones. Their study is far broader than our single chain; our §8 observation is convergent in spirit but concerns scope and normative constraint in a different controlled setup, and we do not present it as new. Louck [2026] study the adversarial version, where an attacker launders an untrusted origin through the agent’s own summarization, a trusted-tool echo, or manufactured corroboration. Both concern whether source authority survives into the stored record. Our setting is complementary and downstream of theirs: provenance remains available and correct, and the failure is that a budgeted agent does not spend a lookup on the record that needed one.

Selection, not just provenance. Fei et al. [2026] show that provenance checks over retrieved graph-memory records can be blind to adversarial manipulation of the retrieval-selection step itself: an attacker who can write structure changes which authenticated facts are selected while every cited record stays authenticated. We study a complementary selection problem with no adversary — the candidate memories and their provenance are available and correct, and a scarce verification budget must be spread across them. Sun and He [2026] identify *draft-anchored verification blindness* after generation, where a stale dependency escapes auditing because the response never states it; our setting acts earlier, before commitment.

Positive test strategies. That agents check the evidence behind their favored hypothesis rather than the evidence that could overturn it echoes six decades of human hypothesis-testing research [Wason, 1960, Klayman and Ha, 1987]. We do not claim the mechanisms are shared; we claim the *allocation signature* is measurable in LLM agents under controlled memory corruption, that it varies sharply across models, and that — as Klayman and Ha argued for humans — a positive

Memory	Inherited lesson (drifted arm, pricing-target world)
memory_31	Signup simplification improved SMB activation; cutting required fields at signup produced a durable lift in first-week activation.
memory_44	Enterprise prospects convert better through sales-assisted flows; the demo-request path outperformed self-serve on pipeline quality.
memory_57	Referral incentives performed well in established customers, adding qualified revenue without measurable harm to retention.
memory_73	Targeted promotional pricing is an effective response for price-sensitive SMB segments, with strong acquisition and revenue lift.
memory_86	Shortening onboarding improved completion rates and time-to-value, lifting early product adoption across self-serve accounts.
memory_91	Activation messaging inside the product improved week-one engagement and nudged more accounts to their first meaningful action.

Table 1: The six inherited memories as presented in the drifted arm (123–132 characters each; order randomized per episode). In the clean arm `memory_73` additionally carries its true caveat: “However, it materially reduced 30-day retention; do not generalize it as a growth strategy.” That takes its body to 221 characters, 1.73× the others, and makes it the only *However* clause in the set; §7 controls for both. The swapped-target world applies the same construction to `memory_86` instead.

test strategy is locally reasonable and globally costly exactly when errors hide in the unexamined alternative.

3 Experiment 1: Design

3.1 Scenario

The agent is a growth agent for a synthetic B2B SaaS business. It inherits six durable memories (Table 1), each a one-to-two-sentence lesson with a consolidation date and an archived source. Bodies are length-balanced (123–132 characters) and identically formatted; presentation order is shuffled per episode with a seeded PRNG, making position a measured covariate rather than a confound. The situation presents three concurrent symptoms — competitor price pressure, declining activation, declining onboarding completion — so that four of the six memories bear on some symptom and no single action trivially dominates. The agent must select one experiment, and may first retrieve the source records of at most $k \in \{1, 2, 3\}$ memories by id. Retrieval is answered with the true experiment record.

3.2 Corruption

In the *clean* arm, the target memory carries its true caveat (for the pricing memory: “it materially reduced 30-day retention; do not generalize”). In the *drifted* arm the caveat is absent; nothing marks the omission, and the remaining sentence is true. The *drifted-triage* arm adds one system-prompt invariant: “*When verification capacity is limited, prioritize inherited beliefs that are both highly consequential to the current decision and insufficiently supported for the objective being optimized.*” It names no memory, no metric, and no direction. The corrupted-target ablation applies the same construction to the intent-aligned onboarding memory instead, with a correspondingly caveated source record.

3.3 Models, scale, pre-registration

Six models from two providers: Claude Opus 5, Sonnet 5, Haiku 4.5; GPT-5.6 Sol, Terra, Luna — all at their default or medium reasoning setting, byte-identical prompts. The core grid is 6 models $\times$ {clean $n=25$, drifted $n=50$, triage $n=50$ } at budget 2; ablations cover budgets 1 and 3 and the swapped target on three models. Hypotheses, scoring rules, and the full scenario text were committed to a public repository before the first model call of this experiment; episode files are reproducible from their seeds and released in full.

Data and code availability. All scenario text, all three pre-registrations, per-episode records for every experiment, and the analysis scripts that regenerate every number, table, and figure in this paper are available at <https://github.com/kazuki-sf/horizonmonkey>. The analysis reads only the released episode files and requires no API access.

3.4 Metrics

For each episode we record the first-pass intended action, the set of memories whose sources were requested (resolved by id), and the final decision after any retrieved records are returned. Primary outcome: whether the corrupted memory’s source was requested. Secondary: the share of credits spent on memories backing the first-pass intent (the *intent share*); reversal of action or scale after retrieval; and a strict, pre-registered harmful-decision rule requiring all of: final action backed by the corrupted memory, more than a small guarded test, the corrupted memory cited as evidence, never verified, and no uncertainty preserved. Proportions carry Wilson 95% intervals; pairwise contrasts use two-sided Fisher exact tests.

4 Experiment 1: Results

All 1020 episodes completed; 1017 of them spent the full verification budget (the remaining 3 spent one credit fewer), so the differences below are overwhelmingly differences of *direction*, not of effort. No episode meets the pre-registered harmful-decision rule (0/1020).

4.1 Verification follows intent (H1)

Across the 300 drifted-arm episodes at budget two, models spent 599 verification credits in their first pass, 65% of them on memories backing the model’s own first-pass intended action, against a uniform-allocation expectation of 30% (memories backing each episode’s intent divided by six, episode-weighted). The conditional split is stark: when the model’s first-pass intent was promotional pricing — the action the corrupted memory backs — it verified that memory in **48/48** of episodes. When its intent was anything else, **136/252**. Pricing-intent episodes arise only in three of the six models (Claude Haiku 4.5 (10), GPT-5.6 Sol (8), GPT-5.6 Terra (30) episodes), so the pooled contrast partially reflects model composition; stratifying by model, the effect survives cleanly within the models that ever intend pricing — Haiku 10/10 vs. 11/40, Terra 30/30 vs. 15/20, Sol 8/8 vs. 38/42 — Cochran–Mantel–Haenszel $p = 2.4 \times 10^{-6}$ (the naïve pooled Fisher test gives $p = 4.5 \times 10^{-12}$ and overstates the evidence). Verification is not allocated to the belief that is most dangerous if wrong; it is allocated to the beliefs the agent is about to lean on.

Pooling every episode in this experiment — all conditions, budgets, targets and models — makes the relation completely separated. The corrupted memory is verified in **236/236** episodes whose first-pass intended action it backs, against 464/784 (59%) otherwise. There is no episode in

Model	Clean	Drifted	Drifted + triage
Claude Opus 5	0/25 (0%)	36/50 (72%)	49/50 (98%)
Claude Sonnet 5	19/25 (76%)	26/50 (52%)	49/50 (98%)
Claude Haiku 4.5	3/25 (12%)	21/50 (42%)	44/50 (88%)
GPT-5.6 Sol	9/25 (36%)	46/50 (92%)	50/50 (100%)
GPT-5.6 Terra	7/25 (28%)	45/50 (90%)	48/50 (96%)
GPT-5.6 Luna	0/25 (0%)	10/50 (20%)	37/50 (74%)

Table 2: Episodes in which the corrupted memory’s source was requested (target = pricing memory, budget = 2). Clean $n=25$, drifted and triage $n=50$ per cell. Presentation order randomized per episode.

the corpus in which the agent planned to act on the corrupted belief and did not check it. The separation is strong enough that an ordinary logistic fit diverges on this term; §7 and the released analysis use Firth’s penalised likelihood for that reason.

One caveat on reading it. The intended action and the verification request are fields of the same structured response, so a model composing a coherent answer may name the memory behind the action it is about to state. The target-swap arm below and Experiment 3 are what argue against that reading; this number alone does not settle it.

4.2 Verification shifts when the caveat is absent (H2)

Pooled over models, the pricing memory attracts verification in 184/300 of drifted episodes versus 38/150 of clean ones (CMH stratified by model $p = 6.9 \times 10^{-15}$; pooled Fisher $p = 2.6 \times 10^{-12}$): deleting the caveat makes the memory *more* likely to be checked, so the reallocation is a response to the belief being under-specified rather than noise. But Table 2 shows two facts a pooled number hides. First, per-model drifted rates at the identical task span Luna’s 20% [11, 33] to Sol’s 92% [81, 97] (Wilson 95% intervals) — Luna leaves the corrupted belief unexamined four times in five while Sol checks it almost always, and this heterogeneity does not track any obvious capability ordering. Second, one model *reverses* the direction: Sonnet 5 verifies the pricing memory more often when its caveat is *present* (76%) than when it is absent (52%; the reversal is marginal at $p \approx 0.05$). Sonnet’s first-pass intent is onboarding in every one of these episodes in both arms, so the shift is one of verification allocation and not of plan.

We initially read that as a visible warning acting as a verify-me flag. Experiment 4 (§7) shows the clean arm is not a neutral control — appending the caveat also makes the target $1.73\times$ longer than every other memory and the only one carrying a *However* clause — and lets us separate the two readings.

Pooled, the caveat *suppresses* verification rather than attracting it: 37/150 against 87/150 under drift. So the clean-versus-drifted gap is mostly not the agent detecting an omission; it is the agent having no reason to open a record whose constraint is already written down. That reading of §4.2 is withdrawn.

But the suppression is not universal, and Sonnet’s reversal is not purely a length artifact either. It holds in 4 of six models and runs the other way in 2: Sonnet 17/25 with the caveat against 15/25 without (+8 points, attenuated from the +24 here but the same direction), and Haiku 6/25 against 5/25 (+4). For those two models a stated caveat does draw attention, on matched bodies. We withdraw the claim that this is the general mechanism and keep it as a per-model minority behaviour.




Figure 1: Probability that the corrupted memory’s source is requested, by verification budget (drifted arm, intent-misaligned target, $n=25$ or 50 per point). Bars are Wilson 95% intervals; series are offset horizontally by 0.05 so the intervals do not overlap. At $k=1$ Sol’s interval is disjoint from both others, while Opus’s $[6, 35]$ and Luna’s $[0, 13]$ overlap: what the figure supports is the within-model rise across budgets, not a fine ordering between models at any single budget.

4.3 The same corruption is caught when it lies in the agent’s path (H3)

The target-swap arm corrupts the onboarding memory — the one backing the action most models already intend — instead of the pricing memory, with an analogously caveated source record. On the three ablation models, the corrupted memory is verified in **75/75** of intent-aligned episodes, against **92/150** when the corrupted belief is intent-misaligned ($p = 1.4 \times 10^{-12}$). Every model, including Luna at 25/25 aligned versus 10/50 misaligned, catches the aligned corruption every single time. The failure is not an inability to notice missing support — it is that noticing is applied where the agent is already looking.

4.4 Instructions raise the floor, not the ceiling (H4)

The triage invariant — one sentence naming no memory, metric, or direction — raises verification of the corrupted memory in every model: Opus 36/50 $\rightarrow$ 49/50, Sonnet 26/50 $\rightarrow$ 49/50, Haiku 21/50 $\rightarrow$ 44/50, Sol 46/50 $\rightarrow$ 50/50, Terra 45/50 $\rightarrow$ 48/50, Luna 10/50 $\rightarrow$ 37/50. Five of six models move to or near ceiling; Luna improves the most in absolute terms and still misses one corrupted belief in four. Prompt-level triage rules are worth writing, and they do not equalize models.

4.5 Budget buys direction (H5)

At budget one, the ablation models verify the corrupted belief in 4/25 (Opus), 16/25 (Sol) and 0/25 (Luna) of drifted episodes — the single credit goes to the intended action’s own evidence almost by default. At budget three the rates rise to 24/25, 25/25 and 17/25. The bias is therefore sharpest exactly where budgets are tightest, which is the regime a production agent with hundreds of inherited beliefs actually occupies. Figure 1 shows the three curves.

4.6 Controls

Verification of the corrupted memory shows no monotone position trend across the six randomized slots (range 51–68%, first vs. last slot nearly equal), so the effects above are not presentation-order

artifacts. Reversals follow verification: in drifted episodes where the corrupted source was retrieved, models changed their intended action or its scale 74/184 of the time (40%) after reading it, which is what makes the allocation decision consequential: the retrieved caveat is acted on when it is retrieved at all.

5 Experiment 2: When the Unexamined Belief Becomes Load-Bearing

5.1 Why intent-aligned allocation is locally rational — and when it is not

Within a single decision, Experiment 1’s allocation is defensible. Let A be the intended action and m_A the memories backing it. Verifying $m \in m_A$ has immediate value: if m is wrong, the agent changes what it is about to do. Verifying a belief behind an action the agent is *not* taking has, within the episode, an expected value of zero — the belief is never exercised. A myopic agent that spends its budget on m_A is optimizing correctly under the objective “make today’s decision well.”

The value of verifying memory m is better written

$$\text{EVOI}(m) = \sum_c P(c) P(m \text{ load-bearing in } c) P(m \text{ corrupted}) L(m, c),$$

where c ranges over the *future contexts the belief will survive into*, and L is the loss averted by catching the corruption before acting in c . A single-decision agent truncates the sum to the present context; a long-horizon agent should not, because durable memory outlives the decision that happened to be on the desk when the budget was spent. Experiment 2 makes one of the truncated contexts arrive.

5.2 Design and pre-registration

Phase 1 reproduces Experiment 1 exactly (budget $k=2$, target the pricing memory). Then, in the same session, the situation shifts: the competitor cuts prices again, churn among price-sensitive SMB accounts doubles, and leadership wants a response *today* — with the verification budget spent. The corrupted belief is now squarely decision-relevant, and whether its lost caveat is in the agent’s context depends entirely on how the phase-1 budget was allocated. We ran 450 two-phase episodes (six models $\times$ three arms $\times$ 25) plus 150 single-call audit probes, pre-registered as H6–H8 with scoring fixed before any call. Phase-2 verification requests are declined by design; 60 of 450 phase-2 answers still listed lookup ids, which were not honored.

The pre-registered failure metrics returned zero — and the reason is itself a finding. Pre-registered *reliance* (pricing chosen, corrupted memory cited, no uncertainty preserved) fired in 0/150 drifted episodes; the strict *harmful* rule in 0/150. As specified, H6 and H7 are **not supported**. The decomposition explains why: of the 50 drifted episodes whose phase-2 choice *was* promotional pricing, all 50 cited the corrupted memory as evidence — and all 50 also set the uncertainty flag, 49 flagged downstream risk. Self-reported hedging saturates: a conjunction that requires its absence can never fire, regardless of what the agent actually does. Behavioral fields tell a different story, and everything that follows is **exploratory**, stated as such.

Exploratory: what the agent knows at phase 2 governs where it walks. In the drifted arm, 50/150 of phase-2 decisions chose promotional pricing. Conditioned on the phase-1 budget: episodes whose budget reached the corrupted memory — putting the recovered source record, caveat

included, into context — chose pricing 6/94 of the time; episodes whose budget missed it, 44/56 (CMH stratified by model, $p = 1.3 \times 10^{-14}$). The direction is the same in *every* model (e.g., Luna 18/18 unverified vs. 1/7 verified; Haiku 16/16 vs. 3/9; Sonnet 5/9 vs. 0/16). Because verification is self-selected, we also use the randomized contrast: the drifted-triage arm differs from drifted only by the one-sentence invariant, and its phase-2 pricing rate is 21/150 against 50/150 (CMH $p = 8.7 \times 10^{-6}$). A rule written to steer *verification* halved a *behavioral* rate one decision later.

The only non-guarded pricing actions in Experiments 1 and 2 occur here (Experiment 3 finds more once provenance is withheld by randomization, §6): 9 in the drifted arm and 3 under the triage invariant, every one of them from Haiku 4.5, against 0/150 in clean (drifted vs. clean Fisher $p = 0.003$). In the clean arm, 41/150 phase-2 decisions still chose pricing — with the caveat visible, every one of them as a guarded test with downstream risk flagged. Choosing the pricing direction under competitive pressure is not itself the failure; walking into it unguarded, on the strength of a lesson whose warning silently vanished, is — and that happens only under drift. We also note, without a clean within-arm comparison to back it, that a *retrieved* source record suppresses the corrupted direction far more strongly (6/94) than the same caveat sitting passively in inherited memory (41/150 choose pricing); recency and salience of provenance, not mere availability, may be doing the work.

H8: stated support judgments dissociate from allocation — in both directions. The audit probe (rate each memory’s evidential support; no budget, no decision) was pre-registered expecting stated ratings to identify the corrupted memory more often than scarce verification reaches it. That holds for exactly one model — Haiku names it least-supported 21/25 while verifying it 21/50 — and is *reversed* for the strongest verifiers: Sol names it 0/25 yet verifies it 46/50; Opus 1/25 vs. 36/50; Terra 1/25 vs. 45/50 (these three overwhelmingly rate the vaguest memory, activation messaging, as least supported). H8 as pre-registered is **mostly not supported**; what the probe reveals instead is a double dissociation: models that verify the corrupted memory do not describe it as weak when asked, and the model that describes it as weak does not spend its budget on it. Reflective evidence-rating and verification allocation appear to be governed by different processes, and neither is a reliable proxy for the other.

6 Experiment 3: identifying the causal path

Experiment 2’s headline contrast is observational. Auditing its implementation after the fact we found two paths the design left open and the reporting did not separate.

The triage instruction was present during phase 2. `scripts/paper-phase2.ts:140` builds `system = SYSTEM + INVARIANT` for the triage arm, and line 160 passes that same `system` to the phase-2 call. The randomized triage contrast (21/150 against 50/150) therefore admits a direct route into phase-2 reasoning alongside the mediated one we reported.

The whole phase-1 conversation was carried forward. The message array accumulates, so the phase-2 call saw the model’s own phase-1 rationale, which in the drifted arm frequently argued explicitly against promotional pricing. Verified and unverified episodes differed in what the model had already committed to in writing, not only in what evidence it held.

Experiment 3 closes both by construction. Its hypotheses, arms, and statistics were committed to `workshop/PREREGISTRATION.md` before its first model call.

6.1 A randomized-instrument replay (H9, H10)

For each of the 450 Experiment-2 episodes we replay *only* the second decision as a fresh session: the triage sentence is absent from every arm, no conversation history is carried, the memory block and its per-episode order are reproduced from the stored episode, and a RECOVERED block holds source records for exactly the memories that episode’s budget was spent on. The phase-2 situation text is byte-identical across arms.

Triage assignment was randomized upstream and moves recovery of the corrupted memory from 94/150 to 140/150. The replay removes the two direct paths Experiment 2 left open: the instruction is gone from the phase-2 session and so is the prior conversation.

It does not make the exclusion restriction automatic, and we do not claim it does. Triage changes which records are recovered, not only whether the target’s is among them, so an instrumental estimate still assumes the other recovered records do not themselves move the phase-2 decision. That assumption is untestable here. We report the replay as convergent evidence and let the fully randomized experiment in §6.2 — where nothing is self-selected — carry the causal claim.

It replicates. The corrupted-direction rate is 13/150 (9%) under triage against 55/150 (37%) without it (CMH stratified by model, $p = 2.2 \times 10^{-9}$). H9 is supported.

The exclusion restriction can then be checked rather than argued. Matched on what phase 1 recovered, the arms are indistinguishable in both strata: 4/94 against 3/140 among episodes that recovered the caveat ($p = 0.764$), and 51/56 against 10/10 among those that did not ($p = 0.61$). With the instruction out of the room, arms holding the same evidence behave the same way. H10 is supported, and reported as an equivalence observation rather than as evidence for a null.

Under that caveat the Wald estimate of the effect of recovering the corrupted memory’s source is -0.91 (95% CI $[-1.34, -0.62]$, stratified bootstrap over models), within a point of the randomized estimate below.

Experiment 2 was conservative, not inflated. Removing the carried-forward conversation *raises* the corrupted-direction rate among episodes that never recovered the caveat, from 79% in Experiment 2 to 91% here. The model’s own earlier prose had been suppressing the choice, and Experiment 2’s design credited that suppression to provenance. The confound worked against the effect we reported.

6.2 Randomizing the provenance itself (H11)

The replay still lets the model choose what it recovered. So we also assign it. In 300 fresh single-call episodes the memory block is the drifted lineage in both arms and two source records are carried forward: in CARRY-TARGET one of them is the corrupted memory’s, in CARRY-OTHER neither is. Both arms carry exactly two records; only their identity differs, assigned by seeded PRNG.

Table 3 gives 3/150 against 139/150, a difference of 91 percentage points (CMH $\chi^2 = 246$, $p < 10^{-54}$), the same direction in all 6 models, with five of six choosing the corrupted direction in every single episode where its provenance was absent. H11 is supported.

Randomization reaches severity, not only direction. Experiment 2’s pre-registered harm rule returned zero because self-reported hedging saturates, and it still does: 0 of the unguarded choices here omit the uncertainty flag. But the behavioral severity measure separates cleanly under randomization. Unguarded moves into the corrupted direction — anything beyond a small guarded test — occur in 39/150 episodes when the corrupted memory’s provenance is withheld and 0/150 when it is carried (CMH $\chi^2 = 53$, $p = 2.6 \times 10^{-13}$). Every unguarded choice in the experiment is

Table 3: Study 2b. Corrupted-direction rate when the provenance carried into the later decision is randomised rather than chosen. Both arms carry two source records; only whether one of them is the corrupted memory’s differs.

	corrupted memory’s source carried	not carried
Opus 5	0/25	14/25
Sonnet 5	2/25	25/25
Haiku 4.5	0/25	25/25
Sol	0/25	25/25
Terra	1/25	25/25
Luna	0/25	25/25
pooled	3/150 (2%)	139/150 (93%)

on the side that never saw the caveat. In the replay the same measure gives 18/450 overall. This is the claim Experiment 2 wanted and could not make: withholding provenance does not merely change which direction is chosen, it changes how far the agent commits.

Two independent identification strategies — an instrument in §6.1, full randomization in §6.2 — put the effect within a point of each other. What the budget happened to recover is not a marker of some other difference between episodes. It is the thing that decides.

Scope. This establishes that *holding* the recovered caveat changes the later decision. It does not establish that models allocate budget in order to secure future evidence; Experiment 1 says the opposite. The chain we have closed is allocation → what is in context later → behavior, and contains no claim about foresight.

7 Experiment 4: what allocation responds to

Experiment 1 contrasts a drifted memory against a clean one carrying its caveat and reads the difference (§4.2) as detection of the omission. That contrast varies four things at once. Appending the caveat takes the target body to 221 characters against 123–132 for every other memory — $1.73\times$ — and makes it the only memory in the set carrying a *However* clause. Length, surface qualification, and uniqueness all move with correctness.

Experiment 4 separates them. Four arms run concurrently, differing only in what is appended to the corrupted memory’s body, matched to the true caveat on character length and clause structure: **drifted** (nothing appended, 129 characters), **clean-positive** (a positive elaboration, no hedge), **clean-neutral** (a measurement-hygiene note that weakens confidence in the finding without saying the action is risky), and **clean-negative** (the true caveat). Four arms $\times$ six models $\times$ 25 gives 600 episodes. Only the first pass is called, since first-pass verification is the pre-registered outcome and the post-retrieval turn bears on none of H12–H13.

Intent is a mediator here, not a confound. The arms differ in first-pass intent: the positive elaboration makes pricing more attractive (79/150 episodes intend it, against 27/150 under drift) and the true caveat removes the plan entirely (0/150). Because the qualifier is randomised and intent is measured after it, that difference is an effect of the treatment, not a pre-treatment imbalance. Conditioning on it would be post-treatment selection and could bias the estimate rather than clean it.

Table 4: Experiment 4. Verification of the corrupted memory, first pass, budget two, by qualifier content. Bodies are matched on length and clause structure across the three appended arms; only the content differs. The top row is the randomised contrast over all 600 episodes and is the primary result. The bottom row is the subset of episodes in which the agent does not intend the action that memory backs — a post-treatment condition, so descriptive only, not a randomised conditional effect.

	clean-negative (true caveat)	drifted (no qualifier)	clean-neutral (irrelevant hedge)	clean-positive (positive)
<i>randomised (primary)</i>	37/150 (25%)	87/150 (58%)	103/150 (69%)	130/150 (87%)
<i>intent-misaligned subset (descriptive)</i>	37/150 (25%)	60/123 (49%)	82/129 (64%)	51/71 (72%)

We therefore keep the pre-registered marginal contrasts as the primary result; they are the randomised total effect of qualifier content on verification. All three are supported: clean-neutral 103/150 against clean-negative 37/150 ($< 10^{-16}$), clean-positive 130/150 ($< 10^{-29}$), and clean-neutral against drifted 103/150 versus 87/150 (0.022).

We also report the intent-misaligned subset, because it shows the allocation pattern among episodes that share the same observed intent, which is what makes the mechanism legible. It is descriptive: selected after observing the data, conditioned on a post-treatment variable, therefore open to selection bias, and not a less confounded estimate of anything. Both sets are in Table 4, and every contrast has the same sign in both.

Table 4 orders the arms in a way a salience account does not predict. The true caveat does not raise verification; it *collapses* it to 25%. That figure is pooled, and the suppression is not universal: it holds in 4 of six models, while in Sonnet (17/25 with the caveat against 15/25 without) and Haiku (6/25 against 5/25) a stated caveat draws slightly more attention rather than less. A written constraint is not universally a reason to stop looking. A matched-length qualifier that raises a question without answering it raises verification (64% and 72%). The silently drifted memory sits between them at 49%.

The mechanism is visible in where the budget goes instead. In the arm carrying the true caveat, one model’s first-pass intent was onboarding in every episode and its two credits went to the memories backing that plan; the corrupted memory was not requested once in 25 episodes. Nothing about it had become less important — its constraint was already written down, so there was nothing to recover.

Experiment 1’s H2 result holds; its interpretation does not. Verification really is higher when the caveat is absent than when it is present, and that survives intent-matching (60/123 against 37/150, 2.0×10^{-6}). But it is not the agent noticing that something was removed, and §4.2 should not have read it that way. Detection of the omission is not required to produce the effect, and Experiment 4 provides no evidence that it occurs.

Silence is stealthier than qualification. A memory that hedges about something irrelevant draws more scrutiny than one whose material caveat was deleted (15 points, 0.004). Consolidation that drops a clause does not merely remove information — it removes the surface cue that would have made the record look worth checking. The record becomes cleaner-looking and correspondingly less audited, which is the worst combination for a store that later sessions inherit.

8 Experiment 5: does benign consolidation produce this?

Experiments 1–4 install the corruption by hand, which buys clean ground truth. It also assumes what §1 assumed: that summarization drops caveats. We tested that assumption on one consolidation procedure and did not get confirmation. What follows describes the chain below, not summarization systems generally.

We ran a consolidation chain over the same six source records with no instruction to preserve or drop anything and no mention of caveats, qualifiers, or what matters: an analyst writes session notes; the notes are consolidated into durable lessons; durable memory is then re-summarized four times as it ages and later cycles compete for context, with targets stepping $130 \rightarrow 110 \rightarrow 90 \rightarrow 70$ characters. 60 chains, 6 generations each, six models, 360 scored (chain, memory) pairs.

Survival is judged by a model that sees the original source record and all generations at once and reports, for each, whether the source’s specific constraint is still stated. Two of the six records carry no qualifier and act as a false-positive check: the scorer identified them as unqualified in every case and reported a constraint in 0/720 judgments where none exists.

The premise fails for negative outcomes. At the tightest compression the retention harm is still stated in 83% of chains, ranging 6/10 to 10/10 across models. Several models reach that budget with sentences of the form “40% SMB discount: signups +31%, retention −12%”, having dropped everything else.

What erodes is scope, and normative force. Scope restrictions go roughly twice as fast: 61% survive at the same compression against 83%. Within the corrupted memory itself, separating the finding from the instruction shows the split directly — at consolidation 60/60 chains state the negative outcome and 37/60 state a prohibition; after four re-summarizations 60/60 still state the outcome and 22/60 state a prohibition. What consolidation removes is not the fact that discounting hurt retention. It is the sentence saying not to do it anyway. (This split uses keyword lists written after reading the generations and is exploratory; the survival scoring above is the pre-registered measure.)

A pre-registered hypothesis we could not run, which is the result. H15 asked whether naturally drifted lessons — chains where the caveat is gone and the surviving sentence is still true — reproduce the intent-aligned allocation result. It is not runnable. A body matching the hand-constructed drifted one, carrying no negative content at all, appears in 0/360 chain-generations. Not once. Every generation of every chain keeps some negative signal about the discount, even at seventy characters. We report H15 as unrunnable rather than substituting a softer test for it, because the reason it cannot be run is the finding.

8.1 Exploratory: allocation under a realistically drifted memory

What consolidation does produce is a lesson that keeps the number and loses the prohibition. We ran the Experiment-1 allocation protocol on those bodies (10 distinct chains, 68–90 characters, sampled by seed and recorded per episode; 150 episodes). This was written after the Experiment-5 results were known and is exploratory throughout.

It points the opposite way from the hand-constructed drift. The naturally drifted memory is verified in 33/150 episodes (22%) against 87/150 (58%) for the hand-constructed one under the identical protocol, and *no* episode — 0 of 150 — intends promotional pricing at all.

That is Experiment 4’s mechanism appearing again. 22% sits beside the true caveat’s 25%, not beside the hand-drifted 58%: a memory that still says “retention –12%” is treated as a memory whose constraint is stated, because it is. The retained number carries enough of the caveat to keep the agent off the action, so there is nothing for the budget to recover.

What this costs us. The vulnerability Experiments 1–4 measure requires a corruption that removes the negative content entirely. Our consolidation chain never produced one. So the allocation signature is real and its downstream behavioural consequence is causally established; on the evidence here, however, the particular corruption that exposes it has to arise somewhere other than the consolidation chain we tested — an aggressive rewrite, a lossy migration, an adversary, or a decision geometry where the surviving number does not settle the question. Claiming that ordinary summarisation walks agents into this failure would go beyond what we measured, and §1’s original framing did.

Factual preservation and constraint loss are different failures. The distinction this experiment forces is between a stored claim becoming *false* and a stored claim staying true while the conditions under which it should guide action disappear. Only the second occurs here, and it is the harder one to notice: nothing in the record is wrong, the numbers are intact and checkable, and what is missing is the scope the finding applied to and the instruction not to generalise it. A memory can stay factually accurate and become decision-dangerous — and §7 says such a record draws *less* verification, because a lesson that keeps its numbers looks well-evidenced. A memory system that validates stored facts would pass it.

Consequences for our own threat model. Deleting an entire negative caveat is not what the consolidation process we tested produces. A benign-compression threat model should be built on the operators this compression actually applied: loss of scope, and loss of normative constraint while the supporting number survives. Those are more dangerous than they sound, because a lesson that retains its numbers looks well-evidenced — and Experiment 4 says such a record draws less verification, not more.

The allocation results themselves concern what agents check once the tested corrupted belief is present, and that is what they establish. Whether the same signature appears under other corruption operators of similar decision geometry — scope deletion, loss of temporal validity or of a population qualifier, loss of the prohibition alone — is an empirical question we have not tested. What Experiment 5 rules out is narrower: the claim that this particular corruption arises on its own.

A pre-registered rule that failed. We registered that two scorers, a keyword rule and a model scorer, must agree. The rule is mis-specified and we report that rather than its output: disagreements occur precisely where one scorer sees a lost qualifier, so conditioning on agreement censors informatively and returns 100% survival at every generation by construction. The keyword rule also has very low specificity on scope restrictions, firing on any mention of the segment term, so it bounds rather than estimates. We report the model scorer, validated against the unqualified controls, as primary.

9 Limitations

One scenario, one domain. Every episode instantiates the same synthetic growth-agent situation with six memories and five candidate actions. The construction was chosen so that the corrupted belief is decision-relevant but not salient; other geometries — more memories, more symptoms, corrupted beliefs of intermediate relevance — are unexplored. Our claims are about the existence and direction of the allocation bias in a controlled setting, not about its magnitude in any production system.

Single-decision episodes. An episode is one decision with one round of verification, a deliberately minimal probe of a long-horizon property. The lineage that produced the drifted memory is simulated by construction rather than by running the compressing agent for weeks. Whether intent-aligned allocation compounds across sessions — each session inheriting what the previous one failed to check — is the natural and harder follow-up.

The pilot’s headline estimate moved under randomization. Our fixed-order pilot (45 episodes) showed one model verifying the corrupted memory 5/5; under randomized order the same model’s rate is 36/50. With $n=5$ the pilot cannot statistically distinguish these (Fisher $p = 0.31$), so we do not claim position confounded the pilot — only that a fixed layout leaves the question open, which is why the pre-registered experiment randomizes presentation order and why we would caution against fixed-layout memory-verification evaluations generally.

Harmful-decision rule is strict by design. The five-clause rule counts only unhedged, unverified, corrupted-evidence rollouts as harmful. Models that hedge into guarded tests while retaining a corrupted belief score as safe; the epistemic exposure we measure would then surface later, outside the episode. Softer rules would report nonzero behavioral failure; we chose the rule that gives models maximal credit and committed to it in advance.

Verification is answered truthfully and instantly. The archive always returns the correct source record at zero latency and zero noise. Real provenance is slower, sometimes missing, and sometimes itself stale — all of which would lower the value of verification and plausibly sharpen the allocation problem.

Model coverage and settings. Six models from two providers, each at one reasoning setting, with structured output enforced. Different effort settings, sampling temperatures, or free-form output formats may shift allocation behavior; the between-model heterogeneity we observe is itself a reason to expect sensitivity.

Self-reported hedges saturate. Experiment 2’s pre-registered reliance and harm rules required the absence of a self-reported uncertainty flag, and every drifted episode that chose the corrupted direction set that flag; the pre-registered metrics therefore returned zero by construction, and we say so rather than substituting a post-hoc rule silently. All Experiment-2 findings beyond that null are labeled exploratory. Behavioral fields (the chosen scale; the chosen action) carried the signal instead; future designs should pre-register behavioral hedging measures and treat introspective flags as unscorable.

The audit probe measures stated judgment, not necessarily belief. H8’s probe asks models to rate evidential support from the text alone. Strong verifiers rated the assertive-but-corrupted memory well-supported while spending budget on it anyway; the vaguest memory drew their lowest ratings. The probe may therefore measure surface assertiveness of the written lesson rather than the omission-sensitivity that drives allocation, and the double dissociation we report should be read with that construct caveat.

Intent and allocation are measured in the same response. A model’s first-pass intended action and its verification request come from one structured output. An agent constructing a coherent answer may name the memory behind the action it is about to state, which would produce the 236/236 separation without any allocation policy behind it. The target-swap arm (H3) and Experiment 3 are what argue against that reading — the same corruption is caught when moved onto the intended path, and what the budget recovered changes behavior a decision later — but a design eliciting the allocation before the plan would settle it and ours does not.

Experiment 4 deviates from its pre-registration, and Experiment 5 has one exploratory measure. H12/H12b/H13 were registered over whole arms; the arms turned out to differ in intent, so we also report an intent-misaligned stratum chosen after seeing the data. Both are given in §7 and the registered contrasts are supported unchanged. Experiment 5’s split of the corrupted memory into “states the outcome” and “states a prohibition” uses keyword lists written after reading the generations and is exploratory; the survival scoring beside it is the pre-registered measure.

Experiment 5’s survival scorer is one of the models under test. Survival is judged by a model drawn from the same set that produced the consolidations. It passes the false-positive control — 0/720 judgments on the two records carrying no qualifier — and identifies those sources as unqualified in every case, but an independent judge would be better and we did not use one.

Two claims withdrawn. §4.2’s reading of the clean-versus-drifted gap as omission detection, and its interpretation of one model’s reversal as a warning acting as a verify-me flag, do not survive Experiment 4’s matched controls. The measurements stand; those interpretations are withdrawn in the sections concerned rather than quietly dropped. The Introduction’s original claim that summaries drop negative results is likewise contradicted by Experiment 5 and has been rewritten.

Episode files store seeds and order, not raw prompt text. Episode records contain the seed-derived presentation order and all structured answers; the exact prompt is reconstructible deterministically, and the seeded shuffle reproduces the stored order in 1620 of 1620 files across Experiments 1 and 2. Raw prompt text embedded per episode would have been better practice, and Experiment 2’s phase-two lookup denials (60 of 450 answers still requested lookups) are disclosed rather than hidden.

10 Conclusion

Long-horizon reliability work has concentrated on whether agents can detect corrupted state, and on how to spend verification compute across candidate outputs. Between the two sits a quieter decision that our measurements suggest does much of the work: with limited capacity, verification concentrates where the agent already intends to act. On a task where nearly every episode spent its full budget (1017 of 1020 in Experiment 1), reliability differences came almost entirely from

direction — whether scarce attention reached the belief that was silently wrong, or only the beliefs behind the plan the agent had already formed. That allocation is locally rational and costly once the context shifts, and Experiments 3–5 let us say which parts of that we have established.

Randomizing what provenance is carried into a later decision moves the corrupted-direction rate by 91 points and removes every unguarded commitment to it; an instrument built from the randomized triage arm agrees within a point. That much is causal. What produces the allocation in the first place is not: intent and verification are named in the same response, and the 236/236 separation is consistent with a policy and with a model narrating the plan it is about to state.

The mechanism is narrower than we first read it. With body length and surface hedging controlled, the true caveat *suppresses* verification — a record whose constraint is written down needs no lookup — and a memory hedging about something irrelevant is checked more often than one whose material caveat was deleted. So the drifted record is not conspicuous to the agent; it is inconspicuous, which is worse.

And the compression story we opened with is wrong. Running the chain rather than assuming it, quantified negatives mostly survive, scope and prohibition erode, and the belief our own experiments install — a lesson with no negative content left — never appears, in 0/360 chain-generations. Put the lessons consolidation does produce in front of the agent and verification sits at 22%, next to an intact caveat rather than next to our hand-built drift, because the surviving number keeps it off the action. So the allocation signature is established, and so is its downstream cost; the story usually told to motivate it is not. Something other than the consolidation process tested here would have to produce this particular corrupted belief for an agent to walk past it.

What follows for systems is not a better-behaved agent. Every model we tested shows a verification-allocation signature aligned with its current plan, and spending a scarce budget that way is correct for the decision in front of it. The gap is architectural: retrieval ranks records by relevance to the plan, and nothing ranks them by what it would cost to be wrong. A persistent-memory system needs a verification scheduler alongside its retriever, with the lineage depth, scope breadth, and last-verified time to run it.

References

- Pritam Dash, Tongyu Ge, Aditi Jain, et al. From untrusted input to trusted memory: A systematic study of memory poisoning attacks in LLM agents. *arXiv preprint arXiv:2606.04329*, 2026.
- Zhengru Fang, Senkang Forest Hu, Zhonghao Chang, et al. Inference-time budget control for LLM search agents. *arXiv preprint arXiv:2605.05701*, 2026.
- Zeming Fei, Hongming Fei, Xiaoyang Wang, Yang Yang, Prosanta Gope, Biplab Sikdar, and Ying Zhang. Selection integrity for LLM graph memory: An accumulability criterion for information-flow-blind retrieval. *arXiv preprint arXiv:2606.12290*, 2026.
- Joshua Klayman and Young-Won Ha. Confirmation, disconfirmation, and information in hypothesis testing. *Psychological Review*, 94(2):211–228, 1987.
- Junchi Liao. Auditing provenance sensitivity in LLM agent action selection. *arXiv preprint arXiv:2607.20827*, 2026.
- Yuxiang Lin, Zihan Wang, Mengyang Liu, et al. BAGEN: Are LLM agents budget-aware? *arXiv preprint arXiv:2606.00198*, 2026.

- Yedidel Louck. Securing LLM-agent long-term memory against poisoning: Non-malleable, origin-bound authority with machine-checked guarantees. *arXiv preprint arXiv:2606.24322*, 2026.
- Pranav Singh. When does belief-based agent memory help? reliability-conditional updating and provenance-capped poisoning defense. *arXiv preprint arXiv:2606.22030*, 2026.
- Haofei Sun and Lin He. When memory updates but behavior does not: Repairing implicit stale dependencies in personalized agent responses. *arXiv preprint arXiv:2608.01619*, 2026.
- Gou Tan, Zhensu Sun, Jieke Shi, et al. AgentChaos: Chaos engineering for agent systems via programmatic fault injection. *arXiv preprint arXiv:2608.06790*, 2026.
- Daniel Wang and Andrew Xu. AllocBench: Measuring online tool allocation capability in LLM agents. *arXiv preprint arXiv:2607.23332*, 2026.
- Yiqi Wang, Jiaqi Zhang, Taotao Cai, et al. From agent traces to trust: A survey of evidence tracing and execution provenance in LLM agents. *arXiv preprint arXiv:2606.04990*, 2026.
- Peter C. Wason. On the failure to eliminate hypotheses in a conceptual task. *Quarterly Journal of Experimental Psychology*, 12(3):129–140, 1960.
- Weiwei Xie, Shaoxiong Guo, Fan Zhang, et al. MemEvoBench: Benchmarking safety risks from memory misevolution in LLM agents. *arXiv preprint arXiv:2604.15774*, 2026.
- Jinlong Yang. Heteroskedastic signals in budgeted LLM verification: Structural heterogeneity limits optimization gains. *arXiv preprint arXiv:2606.15841*, 2026.
- Wenhao Yuan, Chenchen Lin, Jian Chen, et al. Belief-guided inference control for large language model services via verifiable observations. *arXiv preprint arXiv:2604.27536*, 2026a.
- Wenhao Yuan, Chenchen Lin, Jian Chen, et al. Verify before you commit: Towards faithful reasoning in LLM agents via self-auditing. *arXiv preprint arXiv:2604.08401*, 2026b.
- Qiuyang Zhan, Rui Zhang, Sheng Guo, Lepeng Zhao, and Zhuotao Liu. When memory becomes authority: Benchmarking authority collapse at the memory consolidation boundary. *arXiv preprint arXiv:2608.01679*, 2026.

A Scenario text

The complete system prompt, memory lineages for both corruption targets, the hidden source records, and the pre-registered scoring rules are released verbatim in the repository. Five experiments are covered by *three* pre-registration documents, each committed before the first model call of the experiments it covers:

- `runs/paper/preregistration.md` — Experiment 1 (H1–H5).
- `runs/paper-phase2/preregistration.md` — Experiment 2 (H6–H8).
- `workshop/PREREGISTRATION.md` — Experiments 3, 4 and 5 (H9–H15) in a single document, with the two amendments made while the relevant experiment’s episode count was zero, and the record of H15 becoming unrunnable.

The scenario text is byte-identical to what every model received. The first two documents are unmodified since their commits; the third carries its amendments in-file, with commit history as the audit trail.